

Ch 5: ARIMA model

Contents

1	Non-stationarity and Box-Jenkins Method	3
1.1	Non-Stationary Data	4
1.2	Tests for Stationarity	7
1.3	Box-Jenkins Method	15
1.4	Random Trends	21
1.5	Summary 1:	37
2	Examples of ARIMA fitting	38
2.1	ARIMA(p,d,q) model	39
2.2	Lake Huron Data	40
2.3	Sheep Data	43
2.4	Dow Jones Data	45
2.5	Drift vs Constant	51
3	ARIMA(p,d,q) Model Selection	53
3.1	Model Selection in ARIMA(p, d, q)	54
3.2	Overdifferencing	57
3.3	Example: LArain data	65
3.4	Example: Color Data	67

3.5	Example: Global Temp Data	69
4	ARIMA Forecasting	72
4.1	ARIMA forecasting	73

March 22, 2017

Non-stationary Time Series and Box-Jenkins Method


[\[ToC\]](#)

1.1 Non-Stationary Data

[\[ToC\]](#)


Dow Jones Index

From Aug. 28 to Dec. 18, 1972




Lake Huron Data

Level of Lake Huron 1875-1972



Global Temperature Data

Global temperature data from Shumway



1.2 Tests for Stationarity

[\[ToC\]](#)

How do you check if a time series is stationary?

1. Visual inspection for constant mean, constant variance
2. KPSS test (H_0 : Stationary)
3. ADF test (H_0 : Not Stationary)
4. PP test (H_0 : Not Stationary)
5. Fit AR(1), and see if $\hat{\phi}_1$ is significantly different from 1.

KPSS test

Kwiatkowski-Phillips-Schmidt-Shin (1992) test for

- Default choice in `auto.arima()`

$$\begin{cases} H_0 : X_t \text{ is trend stationary} \\ H_A : X_t \text{ is not stationary} \end{cases}$$

- Large p-value means "Stationary".
- Decompose the series as

$$X_t = T_t + W_t + Y_t$$

where T_t is deterministic trend, W_t is random walk, and Y_t is stationary error.

- Lagrange multiplier test that the random walk has zero variance.

ADF test

Augmented Dickey-Fuller test

- "Unit-root" test
-

$$\begin{cases} H_0 : Y_t \text{ is not stationary} \\ H_A : Y_t \text{ is stationary} \end{cases}$$

Is replaced by

$$\begin{cases} H_0 : Y_t \text{ has unit root} \\ H_A : Y_t \text{ does not have unit root} \end{cases}$$

- Small p-value rejects the null hypothesis of non-stationarity, and means "Stationary".

Fit AR(1) to a time series, and test if ϕ_1 is significantly different from 1, by

$$\frac{\hat{\phi}_1 - 1}{SE(\hat{\phi}_1)} \sim \mathcal{N}(0, 1)$$

PP test

Phillips-Perron test for the null hypothesis that x has a unit root.

- Same Hypothesis as ADF test
- Small p-value means "Stationary".
- The tests are similar to ADF tests, but they incorporate an automatic correction to the DF procedure to allow for autocorrelated residuals.
- Calculation of the test statistics is complex.
- The tests usually give the same conclusions as the ADF tests

Three stationarity test

- P-value of (KPSS, ADF, PP)
 - (Large, small, small) → All three indicating stationarity.
 - (Small, large, large) → All three indicating non-stationarity.
 - (Large, Large, large) → Conflicting conclusions, inconclusive.
- `Stationarity.tests()` performs all of the three tests. It is located at the same place as `Randomness.tests()`. Command below will load it into R.

```
source("http://gozips.uakron.edu/~nmimoto/477/TS_R-90.txt")
```

Example:

```
library(forecast)
source("http://gozips.uakron.edu/~nmimoto/477/TS_R-90.txt")

#--- Test it on random sample ---
X <- rnorm(100)
plot(X, type="o")

Stationarity.tests(X)

#--- Lake Huron Data ---
X1 <- read.csv("http://gozips.uakron.edu/~nmimoto/pages/datasets/lake.txt")
X <- ts(X1, start=1875, freq=1)
plot(X, type="o")

Stationarity.tests(X)

#--- Global Temp Data ---
#- install.packages("astsa")

library(astsa)
data(gtemp)
plot(gtemp, type="o", ylab="Global Temperature Deviations")

Stationarity.tests(gtemp)
```

Example:

```
#--- Australian Steel Data
D  <- read.csv("http://gozips.uakron.edu/~nmimoto/pages/datasets/Steel.csv")
D1 <- ts(D[,2], start=c(1956,1), freq=12)
D2 <- window(D1, start=c(1969, 12))

plot(D1, type="o")
Stationarity.tests(D1)

plot(D2, type="o")
Stationarity.tests(D2)

D4 <- window(D2, end=c(1981, 12))
D5 <- window(D2, start=c(1984, 1))

Stationarity.tests(D4)
Stationarity.tests(D5)
```

1.3 Box-Jenkins Method

[\[ToC\]](#)

Backward and Difference Operator

- Define the backward operator B ,

$$BX_t = X_{t-1}.$$

- Then define **difference operator** ∇ (del),

$$\nabla = (1 - B).$$

- For example,

$$\nabla X_t = (1 - B)X_t = X_t - X_{t-1}$$

$$\nabla^2 X_t = (1 - B)(1 - B)X_t = X_t - 2X_{t-1} + X_{t-2}$$

Differencing Linear Trend

- Suppose your time series have a line plus a zero-mean stationary noise.

$$Y_t = a + bt + X_t$$

That means

$$Y_{t-1} = a + b(t-1) + X_{t-1}$$

$$\nabla Y_t = Y_t - Y_{t-1} = b + X_t - X_{t-1}$$

- If X_t is stationary, then $X_t - X_{t-1}$ is also stationary. Thus we have now

$$\nabla Y_t = \mu + M_t$$

We can try to model this with ARMA(p,q) with intercept $\mu = b$.

Differencing Quadratic Trend

If m_t is quadratic, then

$$\begin{aligned}Y_t &= a + bt + ct^2 + X_t \\Y_{t-1} &= a + b(t-1) + c(t-1)^2 + X_{t-1} \\Y_{t-2} &= a + b(t-2) + c(t-2)^2 + X_{t-2}.\end{aligned}$$

That means

$$\begin{aligned}\nabla^2 Y_t &= Y_t - 2Y_{t-1} + Y_{t-2} \\&= 2c + X_t - 2X_{t-1} + X_{t-2}\end{aligned}$$

If X_t is stationary, so is $X_t - 2X_{t-1} + X_{t-2}$.

If m_t is polynomial of deg k , with coefficients c_0, c_1, \dots then applying k power of difference operator will remove the trend.

$$\nabla^k Y_t = k!c_k + \nabla^k X_t.$$

Then you will end up with some stationary series $\nabla^k X_t$ with constant trend $k!c_k$.

Example: Linear Trend

```
t  <- 1:100
Y  <- 3 - .1*t + arima.sim(n=100, list(ar = c(.7,.2)))

plot(Y,type="o")           #- Y is simulated data

Y2 <- diff(Y)              #- take the difference

plot(Y2, type="o" )
```

Example: Quadratic Trend

```
t    <- 1:100
tsq  <- t^2
Y    <- 3 - .5*t + .1*tsq + arima.sim(n=100, list(ar = c(.7,.2)))*10

plot(Y,type="o")

plot(diff(Y),type="o")

plot(diff(diff(Y)),type="o")
```

1.4 Random Trends

[\[ToC\]](#)

Suppose we have a model with trend

$$Y_t = M_t + X_t,$$

where M_t is a random walk:

$$M_t = \sum_{i=1}^t e_i \quad \text{where} \quad e_i \sim_{iid} \mathcal{N}(0, 1).$$

Random Walk

$$M_t = \sum_{i=1}^t e_i \quad \text{where} \quad e_i \sim_{iid} \mathcal{N}(0, \sigma^2).$$

With iid errors (e_1, e_2, e_3, \dots) random walk is generated as

$$\begin{aligned} M_1 &= e_1 \\ M_2 &= e_1 + e_2 \\ M_3 &= e_1 + e_2 + e_3 \\ M_4 &= e_1 + e_2 + e_3 + e_4 \\ &\vdots \end{aligned}$$

We can use `cumsum()` function in R to do this easily.

Mean and Var of RW

$$M_t = \sum_{i=1}^t e_i \quad \text{where} \quad e_i \sim_{iid} \mathcal{N}(0, \sigma^2).$$

Mean is

$$E(M_t) = E\left(\sum_{i=1}^t e_i\right) = \sum_{i=1}^t E(e_i) = 0$$

Variance is

$$V(M_t) = V\left(\sum_{i=1}^t e_i\right) = \sum_{i=1}^t V(e_i) = \sigma^2 t$$

Example: RW

```
e = rnorm(100)      #- 100 random number from N(0,1)
```

```
M = cumsum(e)
```

```
plot(M,type="o")
```

```
plot(M,type="l", ylim=c(-40,40))
```

```
e = rnorm(100)      #- copy and paste these 3 lines many times
```

```
M = cumsum(e)
```

```
lines(M)
```


Random Walk Difference

Since

$$M_t = \sum_{i=1}^t e_i,$$

we have

$$\nabla M_t = M_t - M_{t-1} = \sum_{i=1}^t e_i - \sum_{i=1}^{n-1} e_i = e_t$$

And we know that $e_t \sim \mathcal{N}(0, \sigma)$.

So if M_t is Random Walk, then ∇M_t should look like iid $\mathcal{N}(0, \sigma)$ noise.

Example: RW

```
e = rnorm(100)      #- 100 random number from N(0,1)
M = cumsum(e)
plot(M,type="o")

layout(c(1,2)); plot(diff(M), type="o"); acf(diff(M))
hist(diff(M))
```

Random Walk with a Drift

Random Walk with drift δ is

$$M_t = \sum_{i=1}^t e_i \quad \text{where} \quad e_i \sim_{iid} \mathcal{N}(\delta, 1).$$

That means if M_t is Random Walk with drift, then ∇M_t should look like iid $\mathcal{N}(\delta, \sigma)$ noise.

Example: RW with drift

```
e = rnorm(100, .1, 1)      #- 100 random number from N(.1,1)
M = cumsum(e)
plot(M,type="o")
```

```
plot(M,type="l", ylim=c(-40,40))
```

```
e = rnorm(100, .1, 1)      #- copy and paste these 3 lines many times
M = cumsum(e)
lines(M)
```

Random Walk as Trend

If $Y_t = M_t + X_t$, then

$$\begin{aligned} Y_t &= \sum_{i=1}^t e_i + X_t \\ Y_{t-1} &= \sum_{i=1}^{t-1} e_i + X_{t-1} \end{aligned}$$


That means

$$\nabla Y_t = Y_t - Y_{t-1} = e_t + X_t - X_{t-1}$$

Since X_t, e_t are stationary, so is $e_t + X_t - X_{t-1}$.

Example: Temperature Data

Global temperature data from Shumway



```
library(astsa)
data(gtemp)

plot(gtemp, type="o", ylab="Global Temperature Deviations")


fit <- lm(gtemp~time(gtemp));

layout(c(1,2))
plot(fit$residuals, type="o"); plot(diff(gtemp), type="o")

layout(c(1,2))
acf(fit$residuals); acf(diff(gtemp))
```


Example: Temperature Data

Residuals from regression (Top) vs Differencing with Lag 1 (Bottom)




Example: Temperature Data

ACF of residuals (Top) vs ACF of Differences (Bottom)



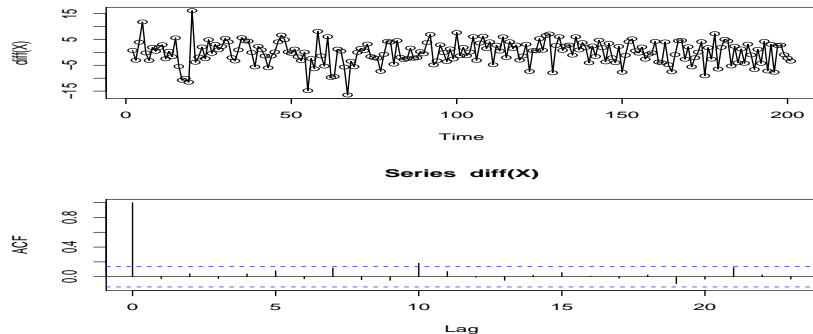
Example: AAPL

Daily adjusted close price of Apple, Inc. Between 10/17/2007 to 8/4/2008 from Yahoo! finance website



Example: AAPL

Difference at lag 1 (Top) and ACF of the difference (Bottom)



```
install.packages("quantmod") #- install if first time

library(quantmod)

getSymbols("AAPL")          #- download from Yahoo!

X <- as.ts(AAPL$AAPL.Adjusted[400:200])
plot(X, type="o")

layout(c(1,2))
plot(diff(X), type="o"); acf(diff(X))
```

1.5 Summary 1:

[ToC]

-
- Three popular test for stationarity is KPSS, ADF, and PP test.
 - Tests are available in R, can be implemented with below command.

```
source("http://gozips.uakron.edu/~nmimoto/477/TS_R-90.txt")  
Stationarity.tests(X)
```

- Classical method tries to identify trend, and tries to remove it by polynomial regression.
- Box-Jenkins Method tries to take a difference between today's data and yesterday's by applying $\nabla = (1 - B)$ operator.
- applying B-J method will often times stationarize non-stationary time series.
- We can use $\text{ARMA}(p, q)$ model to model the stationarized version of the series.
- If you use ∇ once, and apply $\text{ARMA}(2, 3)$, then it is called $\text{ARIMA}(2, 1, 3)$ model.

Examples of ARIMA fitting

[\[ToC\]](#)

2.1 ARIMA(p,d,q) model

[\[ToC\]](#)

-
- Defiend as

$$\nabla^d Y_t = X_t$$

$$\Phi(B) X_t = \Theta(B) e_t$$

where $e_t \sim WN(0, \sigma^2)$.

- So ARIMA(p,d,q) model means you difference the data d times, then you get ARMA(p, q).

2.2 Lake Huron Data

Level of Lake Huron 1875-1972

```
library(forecast)
source("http://gozips.uakron.edu/~nmimoto/477/TS_R-90.txt")

D <- read.csv("http://gozips.uakron.edu/~nmimoto/pages/datasets/lake.txt")
D1 <- ts(D, start=1875, freq=1)

#--- 1. Direct ARMA fit ---
Fit1 <- auto.arima(D1, d=0)    # find best ARMA(p,q) by AICc
Fit1

Randomness.tests(Fit1$resid)

#--- 2. Fit linear trend ---
Reg2 <- lm(D1~time(D1))
summary(Reg2)

plot(D1, type="o"); abline(Reg2, col="red")

Fit2 <- auto.arima(Reg2$residuals, d=0)
Fit2
Randomness.tests(Fit2$resid)
```



```
#--- 3. Direct ARIMA fit ---  
Fit3 <- auto.arima(D1)
```

```
Randomness.tests(Fit3$resid)
```

```
Stationarity.tests(D1)  
plot(forecast(Fit2))  
plot(forecast(Fit3))
```

Lake Huron Data

1.

Y_t is ARMA

2.

$$Y_t = a + bt + X_t \quad X_t \text{ is ARMA}$$

3.

∇Y_t is WN

Y_t is ARIMA(0,1,0)

2.3 Sheep Data

Annual sheep population (1000s) in England and Wales 1867–1939

```
D <- read.csv("http://gozips.uakron.edu/~nmimoto/pages/datasets/sheep.csv", header=T)
D1 <- ts(D[,2], start=c(1867), freq=1)

plot(D1)

#--- 1. Direct ARIMA fit ---
Fit1 <- auto.arima(D1)
Fit1

Randomness.tests(Fit1$resid)

#--- 2. Fit linear trend ---
Reg2 <- lm(D1~time(D1))
summary(Reg2)

plot(D1, type="o"); abline(Reg2, col="red")

Fit2 <- auto.arima(Reg2$residuals, d=0)
Fit2
Randomness.tests(Fit2$resid)
```

Sheep Data

1.

Y_t is ARIMA(2,1,2) without drift

$$Y_t = a + X_t$$


$$\nabla Y_t = X_t - X_{t-1} \text{ is ARMA(2,2) without drift}$$

2.

$$Y_t = a + bt + X_t \quad X_t \text{ is MA(3)}$$

2.4 Dow Jones Data

Dow Jones Index Aug. 28Dec. 18, 1972



```

D <- read.csv("http://gozips.uakron.edu/~nmimoto/pages/datasets/dowj.csv")
X <- ts(D, start=c(1,1), freq=1)
plot(X, type='o')

delX <- diff(X)
plot(delX, type='o')

acf(delX)
pacf(delX)

library(forecast)
Fit1 <- auto.arima(delX, d=0)      #- Fit best ARMA to Y using AIC.
Randomness.tests(Fit1$residuals)
Fit1

#-- Exactly same as below --
Fit2 <- Arima(delX, order=c(1,0,1), include.mean=FALSE) #- same as demean=FALSE
Randomness.tests(Fit2$residuals)
Fit2

```

Current Model

- Differenced once, fit ARMA(1,1) with zero-mean.

$$\nabla Y_t = X_t$$

$$(1 - \phi_1 B) X_t = (1 - \theta_1) e_t$$

where $E(X_t) = 0$.

- This is same as

$$(1 - \phi_1 B) (X_t - X_{t-1}) = e_t - \theta_1 e_{t-1}$$

Constant terms after differencing

- After differencing, data may have non-zero mean.

$$\nabla Y_t = Y_t - Y_{t-1} = \mu + X_t$$

where X_t is zero-mean ARMA.

- Differencing again does remove μ , but it will make X_t to be "over-differenced".
- Instead, we should just model $\mu + X_t$ with non-zero mean ARMA with intercept term.


```
plot(X, type='o')
plot(delX, type='o')

Fit2 <- Arima(delX, order=c(1,0,1), include.mean=FALSE)
Fit2
Randomness.tests(Fit1$residuals)

Fit3 <- Arima(delX, order=c(1,0,1), include.mean=TRUE)
Fit3
Randomness.tests(Fit3$residuals)
```

Current Models

1. $d = 1$, fit ARMA(1,1) with zero-mean.

$$\begin{aligned}\nabla Y_t &= X_t \\ (1 - \phi_1 B) X_t &= (1 - \theta_1) e_t\end{aligned}$$

2. $d = 1$ fit ARMA(1,1) with non-zero mean

$$\begin{aligned}\nabla Y_t &= \mu + X_t \\ (1 - \phi_1 B) X_t &= (1 - \theta_1) e_t\end{aligned}$$

1 was deemed better with AICc.

2.5 Drift vs Constant

[ToC]

-
- If original Y_t has a linear trend, it will show up after differencing.

$$\begin{aligned}Y_t &= a + bt + R_t \\ \nabla Y_t &= Y_t - Y_{t-1} = b + (R_t - R_{t-1}) = b + X_t\end{aligned}$$

- If original Y_t has a Random Walk with drift (cf. p.21), it will also show up after differencing.

$$\begin{aligned}Y_t &= W_t(\delta) + R_t \\ \nabla Y_t &= \delta + (R_t - R_{t-1}) = \delta + X_t\end{aligned}$$

- After differencing, drift term becomes constant.

```
Fit4 <- Arima(X, order=c(1,1,1), include.drift=T)
Fit4

Fit5 <- Arima(delX, order=c(1,0,1), include.drift=F)
Fit5

Fit6 <- Arima(delX, order=c(1,0,1), include.drift=T)
Fit6

# drift in Fit4 and intercept in Fit5 is the same constant
# drift in Fit6 is not.

# Note the numerical difference in AICc in Fit4 vs Fit5
```

ARIMA(p,d,q) Model Selection

[\[ToC\]](#)

3.1 Model Selection in $\text{ARIMA}(p, d, q)$

[\[ToC\]](#)

-
1. Use stationarity test to decide how many times to defference (d).
 2. When $\text{AR}(1)$ parameter is estimated, see if ϕ_1 is significantly different from 1.
 3. Watch for sign of overdifferencing.

auto.arima()

Hyndman-Khandakar algorithm for automatic ARIMA modelling in `auto.arima`

1. The number of differences d is determined using repeated KPSS tests.
2. The values of p and q are then chosen by minimizing the AICc after differencing the data d times. Rather than considering every possible combination of p and q , the algorithm uses a stepwise search to traverse the model space.
 - (a) The best model (with smallest AICc) is selected from the following four:
ARIMA(2,d,2), ARIMA(0,d,0), ARIMA(1,d,0), ARIMA(0,d,1).
If $d = 0$ then the constant c is included; if $d \geq 1$ then the constant c is set to zero. This is called the "current model".
 - (b) Variations on the current model are considered:
vary p and/or q from the current model by ± 1 ; include/exclude c from the current model. The best model considered so far (either the current model, or one of these variations) becomes the new current model.
 - (c) Repeat Step above until no lower AICc can be found.

After `auto.arima()`

- Check parameter significance. If parameters are not significant, then remove from the model.
- For value of d selected by `auto.arima()`, take the difference by hand, and use `Stationarity.tests()` to see what other tests say.
- You can change stationarity test used by `auto.arima()`. Default is KPSS test. Type `?auto.arima` to see syntax of how to change default.

3.2 Overdifferencing

[\[ToC\]](#)

Suppose $\delta X_t = Y_t$

$$Y_t = e_t$$

Take ∇ again, then you get

$$\nabla Y_t = e_t - e_{t-1}$$

Now we got MA(1) with $\theta_1 = 1$. That's not invertible.

Overdifferencing

Suppose

$$\nabla X_t = Y_t$$

$$Y_t = e_t - \theta_1 e_{t-1}$$

If you take ∇ again,

$$\nabla Y_t = (e_t - \theta_1 e_{t-1}) - (e_{t-1} - \theta_1 e_{t-2})$$

$$= (e_t - (1 + \theta_1)e_{t-1} + \theta_1 e_{t-2}).$$

So ∇Y_t is MA(2) with

$$\Theta(z) = 1 - (1 + \theta_1)z + \theta_1 z^2$$

Root is

$$\begin{aligned} \frac{-b \pm \sqrt{b^2 \pm 4ac}}{2a} &= \frac{(1 + \theta_1) \pm \sqrt{(1 + \theta_1)^2 - 4\theta_1}}{2\theta_1} \\ &= \frac{(1 + \theta_1) \pm \sqrt{1 - 2\theta_1 + \theta_1^2}}{2\theta_1} \\ &= \frac{(1 + \theta_1) \pm \sqrt{(1 - \theta_1)^2}}{2\theta_1} \\ &= \frac{2}{2\theta_1} \quad \text{or} \quad \frac{2\theta_1}{2\theta_1} \end{aligned}$$

That's a unit root!

1. Testing Unit-Root in MA(q) polynomials \rightarrow not fully resolved.
2. Test Unit-Root in MA(1) \rightarrow see if $\hat{\theta}_1$ is significantly different from 1.

ARMA with same polynomials

Suppose you have ARMA(1,1) with

$$Y_t - .5Y_{t-1} = e_t - .5e_{t-1}$$

Then this is

$$\begin{aligned}(1 - .5B)Y_t &= (1 - .5B)e_t \\ Y_t &= e_t\end{aligned}$$

So Y_t is just a white noise.

Example: Monthly Oil Price

Cryer p88

```
#--- Load package TSA ---
acf.orig <- acf      # keep the default acf()
library(TSA)
acf <- acf.orig

library(forecast)
source("http://gozips.uakron.edu/~nmimoto/477/TS_R-90.txt")

data(oil.price)
X <- oil.price
plot(X, ylab='Price per Barrel',type='o')

#--- Take difference of lag 1 ---
plot(diff(X),ylab='Change in Log(Price)',type='l')

#--- Take log difference of lag 1 ---
plot(diff(log(X)),ylab='Change in Log(Price)',type='l')

Stationarity.tests(diff(log(X)))
```

```
#--- Look for best ARIMA (this gives us seasonal component) ---  
Fit1 <- auto.arima(log(X))  
Fit1  
  
#--- Look for best ARIMA (supress seasonal component) ---  
Fit2 <- auto.arima(log(X), seasonal=FALSE)  
Fit2  
Randomness.tests(Fit2$residuals)  
  
plot(forecast(Fit2))
```

Forcing some parameter to be 0

```
Arima(diff(log(X)), c(2,0,4), fixed=c(NA,NA,NA,0,0,NA,NA))
```


3.3 Example: LArain data

```
library(forecast)
source("http://gozips.uakron.edu/~nmimoto/477/TS_R-90.txt")

acf.orig <- acf; library(TSA); acf <- acf.orig

data(larain)

plot(larain, type="o")

Arima(larain, order=c(1,0,1))
```

Notes

- When orders of AR and MA matches, watch out for equal value of parameters. It may indicate WN.
- If $MA(1)$ gives you 1, it indicates the over-differencing.

3.4 Example: Color Data

```
library(forecast)
source("http://gozips.uakron.edu/~nmimoto/477/TS_R-90.txt")
acf.orig <- acf
library(TSA)
acf <- acf.orig

data(color)
plot(color, type="o")

Fit1 <- auto.arima(color)

Stationarity.tests(color)
```

Notes

- By default, `auto.arima()` only uses KPSS to decide which d to use. Use *Staionarity.tests()* and plots to make more informed desision.

3.5 Example: Global Temp Data

```
source("http://gozips.uakron.edu/~nmimoto/477/TS_R-90.txt")

D <- read.csv("http://gozips.uakron.edu/~nmimoto/pages/datasets/gtemp.txt")
D1 <- ts(D, start=c(1880), freq=1)
plot(D1, type="o")

#--- fit linear trend inside auto.arima() ---
Fit1 <- auto.arima(D1, xreg=time(D1))
Fit1

Randomness.tests(Fit1$resid)

plot(forecasst(Fit1, xreg=2010:2019, h=10))

#--- fit linear trend by hand ---

Reg1 <- lm(D1~time(D1)); Reg1

Fit2 <- auto.arima(Reg1$resid); Fit2

Randomness.tests(Fit2$resid)

plot(Reg1$resid, type="o")
```

```
Stationarity.tests(Fit1$resid)
```

```
#--- Direct ARIMA fit ---
```

```
Fit3 <- auto.arima(D1)
```

```
Stationarity.tests(diff(D1))
```

```
Randomness.tests(Fit3$resid)
```

Notes

- Parameter in AR(1) should be checked for being ± 1 , which indicates non-stationarity.

ARIMA Forecasting

[\[ToC\]](#)

4.1 ARIMA forecasting

[\[ToC\]](#)

-
- Suppose $\nabla Y_t = X_t \sim \text{ARMA}(p, q)$.
 - Since we know how to forecast ARMA, we know how to get

$$\hat{X}_n(h) = a_1 X_n + \cdots + a_n X_n$$

- How can we calculate $\hat{Y}_n(h)$ so that MSE,

$$E \left[\left(Y_{n+h} - \hat{Y}_n(h) \right)^2 \right]$$

is minimized?

- We have two vectors,

$$(1, Y_0, Y_1, \dots, Y_n), \quad (1, Y_0, X_1, \dots, X_n)$$

- They actually span the same vector space, because

$$\begin{bmatrix} Y_1 - Y_0 = X_1 \\ Y_2 - Y_1 = X_2 \\ Y_3 - Y_2 = X_3 \\ \vdots \\ Y_n - Y_{n-1} = X_n \end{bmatrix} \iff \begin{bmatrix} Y_1 = X_1 + Y_0 \\ Y_2 = X_2 + Y_1 \\ Y_3 = X_3 + Y_2 \\ \vdots \\ Y_n = X_n + Y_{n-1} \end{bmatrix}$$

- Because $Y_{n+1} - Y_n = X_{n+1}$, we can rewrite MSE in Y_t as

$$E\left[\left(Y_{n+1} - \hat{Y}_n(1)\right)^2\right] = E\left[\left((X_{n+1} + Y_n) - \hat{Y}_n(1)\right)^2\right].$$

- Also, note that we can write

$$\begin{aligned}\hat{Y}_n(1) &= a'_0 + a'_1 Y_n + \cdots + a'_n Y_1 \\ &= a_0 + a_1 X_n + \cdots + a_n X_1 + b_0 Y_0 \\ &= \hat{X}_n(1) + b_0 Y_0.\end{aligned}$$

- That means

$$MSE = E\left[\left((X_{n+1} + Y_n) - (\hat{X}_n(1) + b_0 Y_0)\right)^2\right].$$

- If Y_0 is uncorrelated with (X_1, X_2, \dots, X_n) , $b_0 = 0$, it drops out, and we get, given the observations Y_1, \dots, Y_n ,

$$\begin{aligned}
 MSE &= E\left[\left(X_{n+1} - \hat{X}_n(1) + Y_n\right)^2\right] \\
 &= E\left[\left(X_{n+1} - \hat{X}_n(1)\right)^2 + 2Y_n\left(X_{n+1} - \hat{X}_n(1)\right) + Y_n^2\right] \\
 &= E\left[\left(X_{n+1} - \hat{X}_n(1)\right)^2\right] + 2Y_n E\left[X_{n+1} - \hat{X}_n(1)\right] + Y_n^2 \\
 &= E\left[\left(X_{n+1} - \hat{X}_n(1)\right)^2\right] + 0 + c.
 \end{aligned}$$

- We know that $\hat{Y}(1)$ is the minimizer of this.
- Therefore,

$$\hat{Y}_n(1) = Y_n + \hat{X}_n(1)$$

- Similarly,

$$\hat{Y}_n(h) = Y_n + \hat{X}_n(h)$$

- So forecast ARMA(p,q) as usual, then add it to the last observation X_n .